

# Northwestern CS 213 - Spring 2022

## SETI Lab

---

The Search for Extraterrestrial Intelligence (SETI) is a long-standing international effort to try to find alien civilizations through their radio wave (and more recently, optical) emissions. An inhabited planet like ours is as bright as a small star in the radio spectrum due to humanity's numerous radio transmitters, from big AM, FM, shortwave, and TV stations all the way down to the tiny transmitters used by computers and phones to connect to WiFi and cellular access points. We make full use of this spectrum to communicate with one another. One of the basic (and fairly naïve) assumptions of early SETI efforts was that this would hold true for other intelligent civilizations. They may, in fact, communicate using technology built on wild physics that we don't yet understand. But for now, we'll have to see what we can uncover with radio waves.

SETI uses antennas generally employed for radio astronomy to “look” at other stars, recording *signals* over an interval of time. Unlike a regular radio receiver, these receivers do not *tune* to a specific frequency, but rather listen on many frequencies at once—they are *broadband*<sup>1</sup> receivers. They also do not try to decode or *demodulate* the signals they receive at all. The point is to produce huge amounts of signal data. The signal data contains not only possible alien signals, but also many human signals, and a lot of noise from all kinds of natural radio sources (e.g. the supermassive black hole at the center of our galaxy).

The signal data from the radio telescopes is analyzed with computer programs that use *signal processing* algorithms, which sift through the data looking for artificial signals that do not come from our own planet. This is a huge computational (and conceptual) challenge, and a significant breakthrough came with the SETI@HOME project<sup>2</sup>, which allowed ordinary people to volunteer their computers to participate in a large-scale, distributed version of these programs. When it was active, there were over 1.3 million registered users whose computers contribute an average of about two thousand trillion floating point operations per second to the effort (2 *PetaFLOPS*!).

In this lab, you will tackle a greatly simplified and constrained version of this computation, with the goal of trying to make your computer execute the computation as fast as it can<sup>3</sup>. You will write a program that, given a raw signal from a broadband receiver, will try to hunt for signs of intelligence<sup>4</sup> within it, specifically some kind of message. Your program will determine whether a given *frequency band* has unusually high power. In addition, your program will use another program that we provide to determine whether that high-power signal actually appears to come from an *intelligent* source. For fun, we will give you some other programs, an *AM demodulator* and an audio converter, which given the signal and the

---

<sup>1</sup> They're broadband receivers, but we actually search for *narrow-band* signals in the analysis stage

<sup>2</sup> <http://setiathome.berkeley.edu/>

<sup>3</sup> This kind of task, i.e. taking an existing, computationally intensive program, and making it faster with smart programming, is referred to as *performance optimization*.

<sup>4</sup> What would actually constitute *intelligent* information content is a topic of ongoing research!

band, will convert the signal to audio format and let you hear it. The AM demodulator is basically a simple AM radio, and your program will tell us how to tune it.

Of course, this is a systems course, so our focus won't be on the aliens or the signal processing. The goals of the lab are:

- to get you to think about what goes into making a simple parallel algorithm,
- to expose you to low-level parallel programming on a shared memory machine using pthreads,
- and to expose you to the effects of compiler optimization.

The specific algorithm you'll be using also allows you to control how much computation is done per memory read or write, letting you see the effects of the *memory mountain*.

## Contents

|  |    |
|--|----|
| Logistics.....   | 3  |
| Part 1: Detecting Signals .....                          | 5  |
| Part 2: Improving Performance with Parallelization ..... | 8  |
| Understanding the Signal Processing in This Lab .....    | 8  |
| Determining What to Parallelize .....                    | 9  |
| Parallelizing with Threads .....                         | 11 |
| Measuring Time, Performance, and Resource Usage .....    | 12 |
| Suggested Approach .....                                 | 13 |
| Submission .....   | 14 |
| Grading criteria .....                                   | 14 |
| Extra Credit .....                                       | 15 |

## Logistics

This lab can be done in teams of up to two people.

We encourage you to use different computers than you normally might. For example, the computers in the Wilkinson Lab in addition to Moore. Unlike prior labs, this lab will consume A LOT of compute power, and if everyone uses our class servers for testing and debugging, it will be very slow.

To start, copy `~cs213/HANDOUT/setilab-handout.tar` to your protected working directory, and then untar it with `tar xvf setilab-handout.tar`. Inside, you will find the following elements:

- `README` – describes any last minute details
- `Makefile`
- **Generation and Playback Programs**
  - `sigen` – this binary program will generate a signal for you to look at.
  - `amdemod` – this binary program will demodulate the signal and create an audio signal in a binary format.
  - `bin2wav` – this binary program will translate the audio signal to a WAV file, which can be played back with any media player.
  - `analyze_signal` – this program looks for “information content” of your signal, i.e. language.
- **Signal and Filtering Code**
  - `signal.[ch]` – load and store signal files in several ways.
  - `filter.[ch]` – generate filters and apply them to signals (sequential only).
  - `timing.[ch]` – find out how long things take using various methods.
- **`band_scan.c` - a sequential version of the program you’ll write.**
- **Pthread examples**
  - `pthread-ex.c` - how to create and use threads on Unix (and most platforms).
  - `parallel-sum-ex.c` – how to compute the sum of an array in parallel.
- `seti-eval` – allows you to evaluate your program and compare it to others via the web site.

In this lab, you will develop the program `p_band_scan.c`, which is a parallel version of `band_scan.c`. `p_band_scan.c` should behave identically to `band_scan.c`, except that it will allow you to vary the number of processors used to compute the solution, making the time to solution faster. Your `p_band_scan.c` must be correct and it should be performance competitive with our implementation.

You will target performance on the Amdahl servers for this instance of the lab. Amdahl-1..4 are new machines purchased to help teach about parallelism. Each one contains a processor that has a large number of hardware threads. We will dedicate some of these machines for student use, and part for running submitted labs.

Moore is configured to make it possible for you to ssh in (or use FastX) from anywhere. For the Amdahl servers (as well as most other teaching servers/machines), access is restricted to Northwestern networks for security reasons. In order to access the Amdahl servers from non-Northwestern networks, you essentially have two options. First, you can ssh into Moore and then ssh from Moore to an Amdahl machine. The second alternative is to be on the Northwestern VPN and then ssh directly to an Amdahl machine. The Amdahl machines do not support FastX.

The two servers you may use for testing are:

`amdahl-1.wot.eecs.northwestern.edu`

`amdahl-2.wot.eecs.northwestern.edu`

Amdahl 3 and 4 are reserved for grading. Do NOT use those machines.

Your files, accounts, etc, are identical between Moore, Amdahl, and other teaching machines. So you can, for example, develop on Moore, and run on an Amdahl machine.

## Part 1: Detecting Signals

Your first task is to get some signal data from our server. If this were actually SETI@Home, this would be real radio telescope data made available via a project server. For this lab, we've set up a system that generates virtual signals, akin to what you might see if you pointed a radio telescope out into space.

Where SETI@Home looks for signals in a 2.5 MHz<sup>5</sup> frequency band, we are giving you signals in a 200 KHz frequency band to make things more tractable. The signals comprise roughly 30 seconds of data. Typically, these radio signals would be subject to the Earth's rotation. However, to simplify things for you, we will ignore this effect. You can assume that your signal was recorded while the radio telescope was stationary and pointing at a fixed patch of sky.<sup>6</sup>

To get your signal, run the `siggen` binary we provide:

```
Moore> ./siggen -n netid1-netid2
```

Be careful how you type that! It's two netIDs with a dash between them (not an underscore). You'll later need to submit your work using exactly the same netIDs and format.

This will take some time to run, and it will give you several, custom virtual signals. Only one of these will contain an alien message. Other signals may just contain local interference or old radio transmissions. Each signal will be in `'sig-N.bin'`, where N will go from 0 to however many signals we decide to give you. You will also receive a file named `'secret'` that contains a secret key that will be used during submission for authenticating you to our server.

You may modify any of the files in the handout directory, except the generation and playback programs, and you can add new files as needed. You can compile `band_scan.c`, and then use it to test out how things work:

```
Moore> make
Moore> ./band_scan
usage: band_scan text|bin|mmap signal_file Fs filter_order num_bands
```

Here, the `"text|bin|mmap"` argument indicates how the `signal_file` is to be interpreted. `Fs`, `filter_order`, and `num_bands` relate to the problem, and will be given to you. We will describe them a bit later. For the purposes of this lab, you can assume that `Fs` will always be 400,000 (400 KHz).

If you run `band_scan` on a signal that has an alien, it will say something like this:

```
Moore> ./band_scan bin sig.bin 400000 32 10
```

```
type:      Binary
file:      sig.bin
```

---

<sup>5</sup> The choice of this frequency band is not arbitrary. It's centered around the *Hydrogen line*.

<sup>6</sup> Some astronomers now think this may actually be the right way to go—more telescopes, each dedicated to its own, fixed patch of the sky. The energy required to transmit a signal of significant power over several light years may be too much for an “always on” beacon that a directional sweep is intended to pick up.

```

Fs:      400000.000000 Hz
order:   32
bands:   10
Load or map file
Read 11999489 samples
signal average power:      0.336757
 0      0.000100 to      19999.999900 Hz:      0.002624 ** (meh)
 1      20000.000100 to      39999.999900 Hz:      0.000108 * (meh)
 2      40000.000100 to      59999.999900 Hz:      0.000060* (meh)
 3      60000.000100 to      79999.999900 Hz:      0.000061 * (meh)
 4      80000.000100 to      99999.999900 Hz:      0.081506 ***** (NOW)
 5     100000.000100 to     119999.999900 Hz:      0.081501 ***** (NOW)
 6     120000.000100 to     139999.999900 Hz:      0.000042 * (meh)
 7     140000.000100 to     159999.999900 Hz:      0.000023 * (meh)
 8     160000.000100 to     179999.999900 Hz:      0.000017 * (meh)
 9     180000.000100 to     199999.999900 Hz:      0.000019 * (meh)

```

<detailed timing information>

Analysis took 10.719594 seconds

POSSIBLE ALIENS 80000.000100-120000.000100 HZ (CENTER 100000.000000 HZ)

This result means that there is an unusual amount of power centered around the radio frequency 100,000 Hz (100 KHz) in this signal. This may be an alien! If you run against the example file `~cs213/HANDOUT/noalien.sig`, you'll see what a more boring signal looks like. If you run against the example file `~cs213/HANDOUT/alien_msg.sig`, you'll see what an alien signal looks like.

**NOTE: We've checked with our scientists, and they tell us that, for the purpose of this lab, aliens will almost certainly be transmitting on frequencies between 50KHz and 150KHz, so if you see some power outside of this band it, it is likely noise or interference and is safe to ignore.**

If you've detected a large amount of power in a narrow band, you might just have an alien message. To further verify this, you can see if it has something that might be akin to language by using this program:

```
Moore> ./analyze_signal -p sig.bin
```

The `-p` option tells it to print out high-powered pulse lengths it found in the signal, and is optional. Do you notice anything special about the pulse lengths you find? If it found something very interesting, it will output the following:

```
The information content of this signal appears to be very high!
```

At this point we can pretty much assume we have an alien message on our hands. (Again, this project is a little simpler than SETI's task!) Make sure to take note of which signal this is. I.e. if this was `sig-0.bin` and it had a message, you will need to later provide 0 to the `handin` program.

You can then decode the signal for audio playback:

```
Moore> ./amdemod -c 100000 sig.bin
```

Here, the "100000" is the center of the "possible aliens" band you detected using `band_scan`. `amdemod` will produce the output file `sig.out`. You can then convert `sig.out` to a WAV file:

```
Moore> ./bin2wav sig.out
```

This will produce the file `out.wav`, which you can play to hear what your alien sounds like. You'll have to download it to your computer first, but then you can play it with any audio software.

## Part 2: Improving Performance with Parallelization

Now that you've tried out scanning for aliens, let's improve the program so it runs much faster. One of the inherent issues is that scanning multiple bands takes time. Luckily, we have two points in our favor: first that the bands can be processed independently, and second that we have computers capable of running many threads simultaneously.

Your main task for this lab is to create a parallel implementation of `band_scan`, which allows the work to be divided among a number of threads and processors. Your parallel implementation of `band_scan` will look like this:

```
Moore> ./p_band_scan
usage: p_band_scan text|bin|mmap signal_file Fs filter_order num_bands
               num_threads num_processors
```

The two additional arguments denote the number of threads and the number of processors you are to use (processors 0 to `num_processors-1`). You should round-robin your threads over the processors as a starting point.<sup>7</sup>

You need to ensure that `p_band_scan` produces the exact same output as `band_scan` so that we can test and grade it properly. The line of the output ("POSSIBLE ALIENS 80000...") is critical – it is how the testing and grading system finds your answer. **Ensure that your program matches the format that `band_scan` uses exactly.**

You'll want to start with a copy of `band_scan.c`. You'll need to modify the `Makefile` to build it. The remainder of this section explains details of how you will parallelize the existing `band_scan` code.

### Understanding the Signal Processing in This Lab

For the purposes of this lab, a signal is an `N` element array of double precision floating point numbers, where each number indicates the level of the signal (the intensity of the induced current in the antenna) at a given point in time. A signal like this is more specifically termed a discrete-time signal (from electrical engineering), and is also known as a time series (from statistics). Our discrete-time signals are periodically sampled, meaning that the time between one sample and the next is fixed. The inverse of this sample time is called the sample rate, and it is the "`Fs`" in the above. To make this more concrete, suppose we have a sequence of numbers like this in the signal file:

```
0.03
0.50
-0.2
-0.9
-0.6
0.3
```

---

<sup>7</sup> Here, "round-robin" means to assign thread `i` to processor `i % num_processors`.



Let's suppose that the first value (0.03) arrived at time 0, and further suppose the sample rate ( $F_s$ ) is 1000 per second (or 1000 Hz, or 1 KHz). This means the second value (0.50) arrived at time  $0 + 1/1000 = 0.001 = 1$  ms, the next (-0.2) at 2 ms, and so on.

A neat fact of signal processing is that any signal can be represented as a sum of shifted, amplified sine waves at specific frequencies (keyword: Fourier). We can think of the signal as being composed of different "amounts" (amplitudes) of these sine waves, which is essentially what the `band_scan` program computes. Each *band* consists of a group of sine waves. For reasons not important here (keyword: Nyquist), the highest frequency of such a sine wave we can find is half the frequency of the sample rate, so 500 Hz in the above example.

Now you can better understand what `band_scan` is doing, since you know the  $F_s$  that was used (400 KHz), and that resulted in being able to look from 0 Hz to 200 KHz. We asked it to look at the sine waves in that range in 10 bins or bands (0 to 20 KHz, 20 to 40 KHz, and so on). It then found a lot of power (meaning high amplitude sines) around 100 KHz.

`band_scan` does its work by applying band pass filters. A band pass filter passes through sine waves that are in a range of frequencies (the band), and rejects others.

If you look at `band_scan.c`, you'll see that for each band of interest, it first generates a filter design. This form of filter is the simplest variety – it looks like another signal in that it's just an array of  $M$  doubles, where  $M = \text{filter\_order} + 1$ . The bigger  $M$  (`filter_order`) is, the better the filter is, but the more expensive it is to use. Unless  $M$  is truly huge, the time spent in generating the filter is negligible.

Once the filter for the band has been generated, `band_scan.c` next applies it to the signal using a small function called `convolve_and_compute_power()`. `convolve_and_compute_power()` is the function that does the heavy lifting (keyword: convolution). You can easily read the function's code to see what it does, but here is the challenge: if the filter is of size  $M$  and the signal of size  $N$ , then `convolve_and_compute_power()` does  $O(N*M)$  work. As `convolve_and_compute_power()` executes, it also sums up the power in the selected band of the output signal, which is an additional  $O(N)$  work. Once all the bands are completed, the band-sums are evaluated using thresholds to see if there is something interesting to be found. This takes negligible work. Therefore, if you are asked to use  $B$  bands, the whole program does  $O(B*N*M)$  work.

## Determining What to Parallelize

Your goal is to do the  $O(B*N*M)$  work in less than  $O(B*N*M)$  time. Ideally, if you have  $P$  processors, you would be able to do it in  $O(B*N*M/P)$  time. You may be able to achieve this in some cases. For different values of  $B$ ,  $N$ ,  $M$ , and  $P$ , you may find there are different challenges to achieving high performance. For example, you may be bottlenecked by I/O, just reading/writing the files, or by the CPU, or by the memory system.

Parallelism requires not just that you have the ability to do multiple things at once, but also that no intrinsic ordering or dependencies in the algorithm are violated. If you violate them, the result may be incorrect. Correct and slow is better than fast and wrong, so be careful. So, you are looking at the work to find out what chunks you can correctly do together.

There are at least three independent ways to parallelize `band_scan` that we can think of, plus the algorithm could be changed (see extra credit). You are welcome to try out any approaches you think might be interesting. **First, we strongly recommend that you try the straightforward approach of executing bands simultaneously. That should be sufficient for achieving the basic learning and performance goals of the lab.**

When we think of the performance of a parallel program, we need to think beyond just the basic run-time for an example. In particular, we are also interested in how the program *scales* with the problem size and with the number of processors. In a perfectly scalable program, we can always double the number of processors and expect the execution time to be cut in half. Very few parallel programs work this way. In fact, if they do, they are usually called *embarrassingly parallel*. A good performance measurement of `band_scan` would look at the execution time as a function of (1) the number of processors (P), (2) the signal size (N), and (3) the size of filter (M). Another useful view is called a *speedup curve*, where we fix the problem size (N and M), and vary the number of processors, plotting time-with-1-processor / time-with-P-processors as a function of P. Finally, we can simply plot time-with-P-processors as a function of P. This is the view we will use for reporting results on the web page, and for grading.

Each Amdahl machine contains an Intel Xeon Phi 7210, which has 64 cores, each with 4 hardware threads, resulting in 256 available hardware threads.<sup>8</sup> In contrast, Moore has two Xeon Gold 6126s, which have 12 cores, each with 2 hardware threads, resulting in 48-way parallelism. However, each Amdahl hardware thread is much slower than a Moore hardware thread.

---

<sup>8</sup> The use of the terms processor, core, hardware thread, software thread (pthread), and process can be a bit confusing. Here is what it means. A machine may have one or more processors. Each processor is a separate chip mounted on the motherboard of the machine. The processors share the main memory system (DRAMs), although each processor can access memory *near* it faster. A processor can have one or more cores. A core is a complete execution unit plus one or more levels of memory cache. Each core can independently fetch, decode, and execute instructions. Usually, the cores of a processor share an L2 or L3 cache. Each core may have one or more hardware threads. A hardware thread (*hyperthread* is what Intel likes to call this) consists of hardware that can fetch and decode instructions. All the hardware threads of a core share the single execution engine of the core. Their purpose is basically to keep that engine busy by feeding it work. The operating system creates the abstraction of software threads, which are the pthreads you will program in this lab. The OS dynamically maps software threads onto hardware threads. You can ask that it maps a software thread to specific hardware thread on a specific core on a specific processor. The OS also creates the abstraction of processes, which contain one or more software threads running in a shared memory space. These processes are accessed by the programmer through `fork/wait` and similar system calls. The OS implements processes using both software threads and virtual memory management, both of which are tightly coupled with the hardware. In addition, some programming languages (e.g., Scheme, some Java implementations, etc) implement another level of threads and processes on top of the operating system supplied software threads (pthreads) and processes.

Again, try parallelizing across bands first. That should get you to the goal for this project. However, two other ways you might find it possible to get better raw performance and/or better speedup include:

- Exploring the impact of different compiler optimizations to increase the performance of the `convolve_and_compute_power()` function. This is probably the lowest-hanging fruit after having parallelized across bands.
- Parallelizing the individual convolutions themselves. If you look at the `convolve_and_compute_power()` function, you'll see that it does  $O(M)$  work for each of the  $N$  output data points. You could work on those data points in parallel instead of, or in addition to working in parallel across the bands.

## Parallelizing with Threads

You will need to partition the work among multiple processors. To do this, you will use threads, which are explained in some detail in your book. In Linux, the threading interface is called *pthread*s. We provide you with some example code explaining how to use *pthread*s.

The file `pthread-ex.c` shows how to use the basic pthread system calls. It is very important that you supply the `-pthread` option to gcc when compiling code that uses pthreads (see the Makefile). The `pthread_create` system call creates a new thread that starts executing in the function you specify. That is, it looks like a function call, but the caller does not wait for the callee to finish! Instead, the caller and callee continue to run simultaneously. The caller can explicitly wait for the callee to finish by using the `pthread_join` system call.

Note that this is quite similar to the `fork` and `wait` system calls discussed in more detail in class and the book. However, while `fork` creates an entirely new process that is an independent clone of the parent process, `pthread_create` creates a new thread of execution (that starts at the callee function) within the current process. The new thread of execution shares all the memory and other state of the current process with the thread that created it, and with all of the other threads in the process. This means they must carefully coordinate access to the shared memory to avoid serious and difficult to track down bugs. While this is extremely challenging, it is outside the scope of this lab. In this lab, your threads only need to read from shared memory (the input signal). Their writes to the output signal should not need to overlap.<sup>9</sup>

You can create as many threads as you want (and that the operating system has memory to track). The operating system will interleave the execution of these threads in time and across the processors available on the system. That is, the operating system can switch from thread to thread on any given processor, and it can move a thread from one processor to another. This scheduling activity happens on the order of every millisecond or so. It does this to “balance the load” and to maintain *fairness* among all

---

<sup>9</sup> In technical terms, a correct implementation of this lab will have “data race freedom”. However, an incorrect implementation may not.

the threads in the system. However, it is sometimes convenient, especially in a parallel program, to directly control which processor a thread runs on. This is known as *processor affinity*. A thread can advise the operating system of the set of processors it would like to be run on. The `pthread-ex.c` example shows how a thread can request that it only run on a specific processor.

The `parallel-sum-ex.c` example shows how to use pthreads to sum up an array of doubles in parallel.

## Measuring Time, Performance, and Resource Usage

The lab code `timing.[ch]` includes three timing tools. The `band_scan.c` code uses them to measure its own activity. You can use them to find out what the bottlenecks are in your code.

The first timing tool, `get_seconds()`, is measuring the passage of real time using the Unix `gettimeofday()` system call. This can be used arbitrarily in any thread since there is exactly one time across the whole system.

The second timing tool, `get_cycle_count()`, is measuring the passage of real time using the processor cycle counter. This is the most accurate measurement, but it's important to note that each core has its own cycle counter, which can make for confusion if a thread migrates from one core to another.

The third timing tool, `get_resources()`, measures resource usage, including time spent using the processor. You can measure the resource usage of your entire process, and of individual threads. This mechanism gives highly detailed information, but the resolution is much lower than the other timing tools.

## Suggested Approach

### ***Don't panic!***

There is a lot here, which can be overwhelming, especially if you've never been given an existing codebase to work on. Part of what we're trying to do is give you the experience of working within an existing, significant codebase. Remember that what you're doing here is writing `p_band_scan.c`. There is a lot of complexity in this overall system, but you are not writing it from scratch – you're just parallelizing this one component of it. Our own reference implementation of `p_band_scan.c` is based on `band_scan.c` and consists of only about 180 additional lines of C. Ask questions and get help.

1. Read and play with the handout code to get a good sense of how things work. Make sure you try both the `alien_msg.sig` and `noalien.sig` signals, and to get your own signals using `siggen`.
2. Read through `band_scan.c` and `filter.c`, focusing specifically on the `analyze_signal()` and `convolve_and_compute_power()` functions, which do the main work.
3. Understand the `pthread-ex.c` and `parallel-sum-ex.c` code. Start by just compiling it and running it, and then look at it more carefully. `parallel-sum-ex.c` shows how to sum up the elements of an array in parallel.
4. Develop a strategy to parallelize the `band_scan` processing. Read the section on this carefully. **The simplest approach will work fine.**
5. Implement your design of `p_band_scan` using pthreads. It must be possible to specify how many threads to use and how many processors to use at run-time. Note that you will need to change the Makefile. You may find that there are comments in the Makefile that describe how to do this.
6. Test your program to make sure it is correct (compare against the sequential `band_scan` program). For testing, you should be able to use any computer that the lab compiles on. You can also use `seti-eval` (which needs to run on moore) to get your result onto the scoreboard. The scoreboard will indicate whether or not your program is correct regardless of its performance.
7. Start evaluating and enhancing your program for performance. You can do this on any machine in the teaching labs (moore, amdahls, Wilk/Tlab, etc) or elsewhere. If you're on a Wilk or Tlab machine, you can run `~cs213/HANDOUT/seti-perf` to do this evaluation. This is a simple script, so you can also copy it and modify it for your own purposes.
8. When your code is correct and performing well, participate in the performance competition using the `seti-eval` script and the web site.<sup>10</sup>

---

<sup>10</sup> It's important to do performance testing independently (on moore, amdahls, etc) machines BEFORE this step. We will dedicate some of the Amdahl machines for evaluations. On these, `seti-eval` will only do evaluation requests one at a time for fairness and accuracy. Hence, it is possible (and highly likely near the due date) for your request to be waiting in line ("queued up") behind other requests.

## Submission

When you are ready to compare your implementation to others, you can use our evaluation script, like this (don't forget to include your secret with the `-k` flag):

```
Moore> ./seti-eval -t team_name -n netid1-netid2 -s ./p_band_scan -k secret -a 0
```

Where 0 is the signal number corresponding to the alien you found. It may take some time for this to complete. `seti-eval` will work correctly only on moore. This will update the following dashboard:

<http://moore.wot.eecs.northwestern.edu/~cs213/setilab-dashboard.html>

Only team names will appear on the dashboard. Your entry will show up with additional information:

- A WAV file you can click on to hear result
- An analysis of roughly how fast your program was, reported as a scalability graph (more information later) that compares your program with our baseline implementation.
- A ranking of your performance compared to us (we are 1.0) and to other participants
- A note of whether you found the alien or not

Note that the evaluation process may take some time—we may run your program with several different signals, filter orders, numbers of bands, numbers of threads, and numbers of processors to produce the report. Furthermore, only one submission may be tested at a time on each of our test servers, so you will wait in a queue along with anyone else testing their code. **This may take a significant amount of time close to the deadline.**

You can submit as many times as you want and try out different implementations to see just how fast you can make this process go.

**You must also submit your files to Canvas.** Run the following command inside your project directory:

```
Moore> tar cf - *.c *.h Makefile* > netid1-netid2.tar
```

Here, `netid1-netid2` should be replaced with with you and your partner's netIDs. Download that tar file onto your own computer and then upload to Canvas. Multiple submissions are allowed, but only the last one will be graded. Only one person in your group needs to submit the tar file.

## Grading criteria

Your lab will be graded based on three criteria:

- 30% Correctness – it should give the same answers as the sequential `band_scan` program, and you should identify the correct signal as containing the alien message.
- 60% Performance – it should perform similarly to our simple parallel implementation. This is the “TAs” curve in the performance graph shown on the dashboard.
- 10% Code quality – we'll take a look at each submission. Your program should be coherent.

For performance grading, we will use the performance graph seen in the dashboard. As long as your curve is within 10% of our simple implementation's curve, you will get full performance credit.

If your implementation is slower, you will lose 10% of the performance points for every 25% it is slower. Please note that our implementation is pretty basic and should therefore be easy to match or beat. It is compiled with optimizations, however.

## Extra Credit

If your implementation is faster than our baseline parallel implementation (whose performance you'll be able to see on the web site), you'll receive extra credit, up to a maximum of 20% more for an implementation that's twice as fast as ours.

To get a large improvement in performance you'll have to investigate other techniques for speeding up the processing. Here are some ideas:

- **Investigate compiler optimizations**, including those that are specific to the machine you are running on (the Amdahls). See if the compiler can improve performance for you. gcc has hundreds of optimization features you can enable/disable from the command line. You can also try clang, which generally does an even better job. Again, there are hundreds of options.
- **Parallelize** `convolve_and_compute_power()`. If you look at this function, you'll see that it does  $O(M)$  work for each of the  $N$  output data points. Can you make it do this work in parallel? One approach is **vectorization**. Both the Amdahls and Moore support the AVX512 vector instructions, allowing for parallelism within a hardware thread. Each core has an AVX512 functional units (24 total on Moore, 64 total on Amdahl). In this function, each AVX512 functional unit can do 8+ operations at a time. You can try out vectorization as part of trying out compiler optimizations in gcc and clang (see above). Vectorization is a compiler optimization. However, the way the function is currently structured makes it very difficult for the compiler to do this optimization. It is possible to restructure it to compute the same result while enabling this optimization.
- **Parallelize using OpenMP**. OpenMP is an extension to C/C++ and other languages that allows shared memory parallelization without explicitly creating threads. It is available in gcc, clang, and many other compilers. Read about OpenMP and parallelize the sequential code in `band_scan.c` using it.
- **Try a different algorithm**. Some say that a Fast Fourier Transform (FFT) can radically change the total amount of work needed to compute what `band_scan` computes, and make the amount of work independent of the number of bands. Try to learn about and implement this approach.